

Reflection – Metaobjects

value-based

Reflection uses constant values of a single built-in type:

```
info x = reflect(argument);
```

Reflection uses constant values of multiple types:

```
auto x = reflect(argument);
```

type-based

The actual (implementation-defined) type might be:

```
template <info X>
struct __metaobject {
    constexpr operator info() const {
        return X;
    }
};
```

Reflection – APIs

value-based

Either purely `constexpr` or template functions with non-type template parameters:

```
constexpr auto foo(info mo);
```

```
template <info MO>  
constexpr auto bar();
```

type-based

`constexpr` template functions taking metaobjects as function arguments:

```
constexpr auto foo(metaobject auto mo);  
constexpr auto bar(metaobject auto mo);
```

```
template <typename T>  
concept metaobject = unspecified;
```

Reflection – Implementation

value-based

Very fast to compile:

```
constexpr auto foo(info mo);
```

Somewhat slower to compile:

```
template <info MO>  
constexpr auto bar();
```

type-based

Very fast to compile, but requires `constexpr` conversion from metaobject to info:

```
constexpr auto foo(info mo);
```

Slower to compile:

```
template <info MO>  
constexpr auto bar(__metaobject<MO> mo);
```


Reflection – Containers

value-based

Spans, vectors, etc. must be fixed to work in `constexpr`.

```
span<info> x = members_of(...);
vector<info> y = bases_of(...);
```

Can be used with some STL algorithms, unless splicing is involved.

type-based

Containers (sequences) are metaobjects themselves.

```
auto x = get_data_members(...);
auto y = get_base_classes(...);
```

Have their own implementation of reflection-related algorithms, splicing is no problem.

Reflection – Pros

value-based

- Faster to compile
- Uses less resources to compile

type-based

- Consistent and unified API
- More friendly to generic programming
- Plays better with ADL
- Better usability
- Easier to teach

Reflection – Cons

value-based

- Inconsistent API
- The `foo(...)` vs. `bar<...>()` syntax makes it less generic
- Rules when to use which, are sort of complicated and may look arbitrary
- More complicated to teach
- Issues with ADL on NTTPs

type-based

- Slower to compile
- Uses more resources to compile

How to materialize 100'000s of metaobjects?

Use a shell script...

```
L=100 # number of repeats
S=1000 # sampling step size
for l in $(seq 1 ${L})
do
  N=$((l * S))
  # factorize N into three integers
  D=...; E=...; F=...
```

...to generate a C++ source file...

```
int main() {
  return bool(qux(make_index_sequence<${D}>{})) ? 0 : 1;
}
```

..., compile and measure:

```
time $(CXX) $(CXXFLAGS) -o /dev/null $<
done
```


The boilerplate – level 1

```
template <size_t ... K>
constexpr auto qux(index_sequence<K...>) {
    return ( ... + baz(
        integral_constant<size_t, K>{},
        make_index_sequence<${E}>{}));
}
```

The boilerplate – level 2

```
template <size_t K, size_t ... J>
constexpr auto baz(
    integral_constant<size_t, K>,
    index_sequence<J...>) {
    return ( ... + bar(
        integral_constant<size_t, K>{},
        integral_constant<size_t, J>{},
        make_index_sequence<${F}>{}));
}
```

The boilerplate – level 3

```
template <size_t K, size_t J, size_t ... I>
constexpr auto bar(
    integral_constant<size_t, K>,
    integral_constant<size_t, J>,
    index_sequence<I...>) {
    // Simulate the metaobject "id" as:
    // MOID =
    //     K * $((N / D)) +
    //     J * $((N / (D * E))) +
    //     I;
    return /* Do something with MOID... */
}
```

The baseline

Just sum the *MOID* values at compile-time

```
template <size_t K, size_t J, size_t ... I>
constexpr auto bar(
    integral_constant<size_t, K>,
    integral_constant<size_t, J>,
    index_sequence<I...>) {
    return (... +
        K * $((N / D)) +
        J * $((N / (D * E))) +
        I);
}
```

Measure how long does this take to compile and subtract from “real” measurements.

Type-based metaobject & template function

```
template <size_t M>
struct wrapper {
    consteval operator size_t() const {
        return M;
    }
};
```

```
template <size_t M>
consteval size_t foo(wrapper<M> w) {
    return w;
}
```

```
return ( ... + foo(wrapper<MOID>{}));
```

Type-based metaobject & consteval function

```
template <size_t M>
struct wrapper {
    consteval operator size_t() const {
        return M;
    }
};

consteval size_t foo(size_t m) {
    return m;
}
```

```
return ( ... + foo(wrapper<MOID>{ }));
```

Value-based metaobject & template function

```
template <size_t M>  
constexpr size_t foo() {  
    return M;  
}
```

```
return ( ... + foo<MOID>());
```

Value-based metaobject & consteval function

```
constexpr size_t foo(size_t i) {  
    return i;  
}
```

```
return ( ... + foo(MOID));
```


Test hardware


- Old desktop¹:
 - i5-2400U @ 3.10GHz (4 cores)
 - 24GB RAM
- Corporate dev laptop²:
 - i7-1185G7 @ 3.0GHz (8 cores)
 - 32GB RAM
- Mid-range gaming laptop³:
 - AMD Ryzen7 4800HS (16 cores)
 - 16GB RAM
- RPi 4B
 - ARM v7l
 - 4GB RAM

¹2010


²2021

³2019

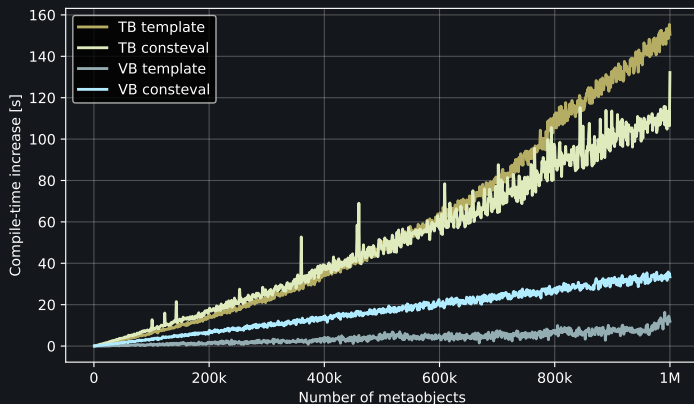
i5-2400 – compile time increase per N metaobjects




i5-2400 – compile time increase per 1 metaobject




i7-1185 – compile time increase per N metaobjects




i7-1185 – compile time increase per 1 metaobject




Ryzen7-4800HS– compile time increase per N metaobjects



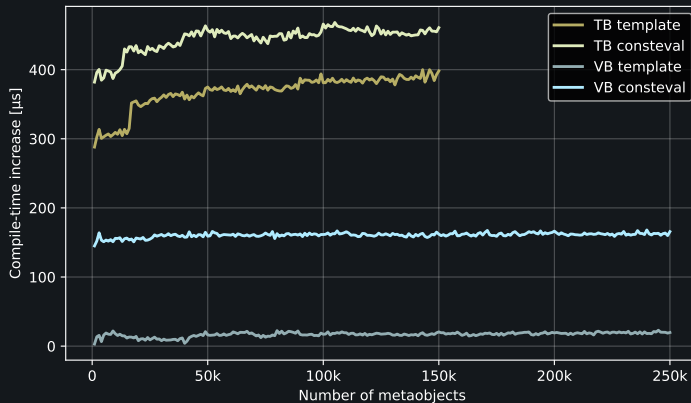
Ryzen7-4800HS – compile time increase per 1 metaobject



ARMv7– compile time increase per N metaobjects



ARMv7– compile time increase per 1 metaobject



The cost of reflection in a “large-ish” project?

- Let's try clang
 - Estimate the number of “things” to reflect
 - Measure the overall compilation time
 - Measure the contribution of reflection
 - Compare purely value-based and typed metaobjects

Estimating number of declarations in clang

Let's try *documented* declarations

Edit doxygen-cfg.in:

```
- GENERATE_XML    = NO
+ GENERATE_XML    = YES
```

Configure:

```
cmake \
  -DLLVM_ENABLE_DOXYGEN=On \
  ...
```

Generate Doxygen docs:

```
ninja doxygen-clang
```

Merge into a single XML file clang.xml:

```
xsltproc combine.xslt index.xml > clang.xml
```

Counting documented declarations in clang

Create count.xslt:

```

<?xml version="1.0" encoding="utf8"?>
<xsl:stylesheet version="1.0"
  xmlns:xsl='http://www.w3.org/1999/XSL/Transform'>
  <xsl:template match="/">
    <xsl:value-of select="count(
      descendant::compounddef4 |
      descendant::member5 |
      descendant::value6 |
      descendant::para7 |
      descendant::param8)" />
  </xsl:template>
</xsl:stylesheet>

```

⁴ struct, class, enum, ...

⁵ data members, member functions, enumerators, ...

⁶ enumerator values, default arguments, ...

⁷ function/constructor/operator parameters, ...

⁸ template parameters, ...

Counting documented declarations in clang

Measure!

```
xsltproc \
    count.xslt \
    clang.xml
```

The result:

379091

- That's for version 15.0.0
- Around FEB-05-2022
- Round that up to 400'000, 500'000 or even 1'000'000
- Let's assume we want to reflect every single declaration

Clean build of clang

Edit `toolchain.cmake`:

```
set(LLVM_USE_LINKER lld)  
set(CMAKE_EXE_LINKER_FLAGS -fuse-ld=${LLVM_USE_LINKER})  
set(CMAKE_SHARED_LINKER_FLAGS -fuse-ld=${LLVM_USE_LINKER})
```

Configure:

```
cmake \  
  -DLLVM_ENABLE_PROJECTS="clang;clang-tools-extra" \  
  -DLLVM_ENABLE_RUNTIMES="libcxx;libcxxabi" \  
  -DLLVM_TOOLCHAIN_FILE="toolchain.cmake" \  
  ...
```

Build and measure elapsed time:

```
time ninja -j 16 install install-cxx install-cxxabi
```


Conclusions

- The typical compile-time overhead of materializing a metaobject is on the order of tens or hundreds of microseconds
- Most typical reflection use-cases don't require reflecting every declaration in a project
- Even if reflecting almost everything, the overhead compared to total build time is a fraction of a percent even in the worst case

The big question

Is the improvement in compile-time worth the decrease in usability of the value-based reflection API?